

Degradable Objects : An Application To A Distributed Data Store

Adam CHADER

PDS track, Master of Computer Science

Institut Polytechnique de Paris

Palaiseau, France

adam.chader@telecom-paris.fr

Abstract—Ceci est un abstract

Index Terms—Concurrent programming, parallel programming, performance, databases

I. INTRODUCTION

This paper is the report of a research project realised for the Parallel And Distributed Systems master at *Institut Polytechnique de Paris*. It describes the work that I realised to apply degradable data structures to Apache Ignite [2], a distributed data store, to improve its performance and scalability.

A. Context

Up until the beginning of the 21st Century, there was a simple rule that could be used to predict the performances of computers. This is known as Moore's Law (Fig. 1). This law states that every two year, the number of transistors would double on chips. This rule could be used quite flawlessly to predict the speed of software, as there is a direct correlation between the speed of sequential programs and the frequency of CPUs. However, in recent years, several issues have surfaced regarding these assumptions : Firstly, we seem to have reached the minimum physical size of a transistor at around 5 nanometer ; secondly, the frequency that can be reached by cramming all these transistors in a single CPU core is too high to be dissipated by current means [4]. As of today, it is extremely complicated and costly to significantly increase the clock speed of a single core. The solution that was found is to put multiple CPU cores in a single machine. The total number of clock cycle per second should continue increasing, while the heat would remain manageable. Intuitively, this solution is perfect : by having two cores instead of one, we could double the total frequency of the CPU, and thus double the speed of software. This is what we call *linear scalability*, and it is still just a dream. There are many reason as to why it is impossible to achieve linear scalability, but they all boil down to the same issue : in order to make the most out of several CPU cores, we need to have them working at the same time. It is possible when the tasks are perfectly parallelisable, but this almost never happens when we want the cores working together. Programs need to access variables in memory, and accessing them in parallel can cause conflicts (we will explain these conflicts in detail later), they also need to aces IO which can cause an issue. The resulting effect of these issues is

that parallel programs often contain sequential portions, where each core has to wait for one core to finish a task. The more sequential portions a parallel program contains, the slower the speedup is. This is described by Amdahl's Law [5] (Fig. 2). It describes the speedup of parallelisation as a function of the amount of sequential code.




Fig. 1: Moore's Law.

```
if(!listPollTime.contains(pollTime.get()))
listPollTime.add(pollTime.get());
PollCounter.add(pollCounter.get());
if(!listPollTime.contains(pollTime.get()))
listPollTime.add(pollTime.get());
Counter.add(pollCounter.get());
```




Fig. 2: Amdahl's Law

With all this in mind, it becomes clear that the burden of increasing the performances of software no longer lies in

the hand of hardware manufacturers, but rather on the hands of software engineers, who should attempt to have as little sequential code as they can in their programs. Historically, it was the responsibility of the end programmer to parallelise their program. But there has been work to hide this work in the lower levels of the software stack [1], to make the shift from sequential to parallel transparent to the higher levels. The scope of this work is at a lower level. The goal is to provide perfectly linearly scalable implementations of classic data structures, keeping the exact same interfaces, such that there could be virtually no modification to applications' source, while achieving the scalability that we desire.

B. Work

This project is a part of a bigger project around *Degradable Data Structures*. These structures will be explained in more detail further in this paper. Simply put, these data structures are created by tweeking the specifications of classical data structures, in order to be able to change their implementations to better ones allowing for better scalability or performance. The goal of this project is to apply these data structures to a specific use case being the data store Apache Ignite [2], and observe the evolution of performances.

C. Contributions

This project made the following contributions to Apache Ignite:

- Benchmarking
- Performance Analysis
- Analysis of the source to find possible applications of degradable objects

D. Outline

The rest of this paper details to context and the contributions of the project.

We will first explain in detail the theory behind degradable objects and how they can be implemented.

We next will detail the information we gathered by analysing Apache Ignite.

we will then show the result obtained.

And finally, we shall compare this to related work.

II. DEGRADABLE DATA STRUCTURES

Introduction STP

A. Contention and Scalability in data structures

zsh:1: command not found: :wq

III. PERFORMANCE ANALYSIS OF APACHE IGNITE

IV. PERFORMANCE EVALUATION

V. RELATED WORK

REFERENCES

- [1] Austin T. Clements, M. Frans Kaashoek, Nickolai Zeldovich, Robert T. Morris, and Eddie Kohler. 2015. The scalable commutativity rule: Designing scalable software for multicore processors. *ACM Trans. Comput. Syst.* 32, 4, Article 10 (January 2015), 47 pages.
- [2] Apache Ignite

- [3] Google Cloud Platform Compute Engine
- [4] R. R. Schaller, "Moore's law: past, present and future," in *IEEE Spectrum*, vol. 34, no. 6, pp. 52-59, June 1997, doi: 10.1109/6.591665.
- [5] Gene M. Amdahl (1989), Oral history interview with Gene M. Amdahl, Charles Babbage Institute, University of Minnesota, hdl:11299/104341.